Transformer with PyTorch

Richard Xu

December 14, 2025

1 Dot-Product Attention (DPA)

1.1 single query $\mathbf{q}$

The very famous paper, "Attention Is All You Need" (Google) [?] introduced the concept of attention, which is a key component of the Transformer architecture

We are given a single query vector $\mathbf{q}$, and matrix of keys $\mathbf{K}$ and values $\mathbf{V}$:

$$\mathbf{q} \equiv \underbrace{\begin{bmatrix} - & \mathbf{q} & - \end{bmatrix}}_{\in \mathbb{R}^{1 \times d_k}} \qquad \mathbf{K} \equiv \underbrace{\begin{bmatrix} - & \mathbf{k}_1 & - \\ - & \dots & - \\ - & \mathbf{k}_m & - \end{bmatrix}}_{\in \mathbb{R}^{m \times d_k}} \qquad \mathbf{V} \equiv \underbrace{\begin{bmatrix} - & \mathbf{v}_1 & - \\ - & \dots & - \\ - & \mathbf{v}_m & - \end{bmatrix}}_{\in \mathbb{R}^{m \times d_v}} \tag{1}$$

and let $(\mathbf{q}, \mathbf{K}, \mathbf{V})$ be tuples. Bear in mind that in Eq.(1), $\mathbf{q}$ is a row vector. This is to conform with Eq.(7), where all the $\mathbf{Q}, \mathbf{K}, \mathbf{V}$ have each element expressed as a row vector.

the Dot-Product Attention (DPA) is defined as:

$$\implies \mathbf{q}\mathbf{K}^\top = \begin{bmatrix} \underbrace{\mathbf{q}\mathbf{k}_1^\top}_{\in \mathbb{R}} & \cdots & \underbrace{\mathbf{q}\mathbf{k}_m^\top}_{\in \mathbb{R}} \end{bmatrix}$$

$$\implies A(\mathbf{q}, \mathbf{K}, \mathbf{V}) \equiv \text{softmax}(\mathbf{q}\mathbf{K}^\top)\mathbf{V}$$

$$= \text{softmax}\left(\begin{bmatrix} \mathbf{q}\mathbf{k}_1^\top & \cdots & \mathbf{q}\mathbf{k}_m^\top \end{bmatrix} \right) \begin{bmatrix} - & \mathbf{v}_1 & - \\ - & \dots & - \\ - & \mathbf{v}_m & - \end{bmatrix} \tag{2}$$

$$= \sum_{i=1}^m \underbrace{\frac{\exp[\mathbf{q}\mathbf{k}_i^\top]}{\sum_j \exp[\mathbf{q}\mathbf{k}_j^\top]}}_{\mathbf{a}_i} \mathbf{v}_i$$

$$\in \mathbb{R}^{d_\mathbf{v}}$$

1.2 Scaled Dot-Product Attention (SDPA)

the problem starts as d_k becomes larger, variance of $\mathbf{q}\mathbf{k}^\top$ increases. hen as a result, some dot product values gets very large, with $\exp(.)$, softmax $\mathbf{p}$ gets peaky! Remember derivative of cross-entropy between softmax p and y is:

$$\mathbf{C}(\mathbf{z}) = -\sum_{k=1}^{K} y_k \left[\log\left(p_k\right)\right] = -\sum_{k=1}^{K} y_k \left[\log\left(\frac{\exp^{z_k}}{\sum_l \exp^{z_l}}\right)\right] \tag{3}$$

$$\implies \frac{\mathbf{C}(\mathbf{z})}{\partial \mathbf{z}} = (\mathbf{p} - \mathbf{y})$$

with a peaky softmax, lots of element in gradient vector $\frac{\mathbf{C}(\mathbf{z})}{\partial \mathbf{z}}$ are zero!
the solution then becomes that of scaling by length of d_k:

$$A(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \mathrm{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_k}}\right)\mathbf{V} \tag{4}$$

"mask (opt.)" is only used at decoder during training

1.3   in the case of "seq2seq with attention"

we have:

$$\mathbf{q} \equiv \mathbf{h}_i$$
$$\mathbf{k}_i = \mathbf{v}_i = \mathbf{z}_i \tag{5}$$
$$\implies A(\mathbf{q}, \mathbf{K}, \mathbf{V}) \equiv A(\mathbf{h}_i, \mathbf{Z}, \mathbf{Z}) = c_i$$

where is our conditional or context vector. In order to confirm the notation in Eq.(1), we have assumed $\mathbf{h}_i$ to be a row vector, so instead of $\mathbf{h}_i^\top \mathbf{z}_j$, we express it as $\mathbf{h}_i \mathbf{z}_j^\top$:

$$a_{ij} = \frac{\exp\left(e_{i,j}\right)}{\sum_{t=1}^{m} \exp\left(e_{i,t}\right)} = \frac{\exp\left(\mathbf{h}_i \mathbf{z}_j^\top\right)}{\sum_{t=1}^{m} \exp\left(\mathbf{h}_i \mathbf{z}_t^\top\right)} \tag{6}$$

2

1.4   multiple queries

now we have multiple $\mathbf{Q} = \{\mathbf{q}_i\}$, e.g., N words in the decoder, we now have $\mathbf{Q}$, $\mathbf{K}$ and $\mathbf{V}$ expressed as:

$$
\mathbf{Q} \equiv \underbrace{\begin{bmatrix} - & \mathbf{q}_1 & - \\ - & \dots & - \\ - & \mathbf{q}_n & - \end{bmatrix}}_{n \times d_k}
\qquad
\mathbf{K} \equiv \underbrace{\begin{bmatrix} - & \mathbf{k}_1 & - \\ - & \dots & - \\ - & \mathbf{k}_m & - \end{bmatrix}}_{m \times d_k}
\qquad
\mathbf{V} \equiv \underbrace{\begin{bmatrix} - & \mathbf{v}_1 & - \\ - & \dots & - \\ - & \mathbf{v}_m & - \end{bmatrix}}_{m \times d_v}
\tag{7}
$$

$A(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \mathrm{softmax}(\mathbf{Q}\mathbf{K}^\top)\mathbf{V}$ can be expressed as follows:

$$
\implies \mathbf{Q}\mathbf{K}^\top = \begin{bmatrix} \mathbf{q}_1\mathbf{k}_1^\top & \dots & \mathbf{q}_1\mathbf{k}_m^\top \\ \dots & \dots & \dots \\ \mathbf{q}_n\mathbf{k}_1^\top & \dots & \mathbf{q}_n\mathbf{k}_m^\top \end{bmatrix}
\tag{8}
$$

in the case of self attention, where $m = n$, then $\mathbf{Q}\mathbf{K}^\top$ is a square matrix of $n \times n$.

$$
\begin{aligned}
\implies \mathbf{O} &\equiv \mathrm{softmax}(\mathbf{Q}\mathbf{K}^\top)\mathbf{V} \\
&= \mathbf{A}\mathbf{V} \\
&= \underbrace{\begin{bmatrix} \mathrm{softmax}\left(\begin{bmatrix} \mathbf{q}_1\mathbf{k}_1^\top & \dots & \mathbf{q}_1\mathbf{k}_m^\top \end{bmatrix} \right) \\ \vdots \\ \mathrm{softmax}\left(\begin{bmatrix} \mathbf{q}_n\mathbf{k}_1^\top & \dots & \mathbf{q}_n\mathbf{k}_m^\top \end{bmatrix} \right) \end{bmatrix}}_{n \times m} \underbrace{\begin{bmatrix} - & \mathbf{v}_1 & - \\ - & \dots & - \\ - & \mathbf{v}_m & - \end{bmatrix}}_{m \times d_v} \\
&= \underbrace{\begin{bmatrix} \underbrace{\sum_{i=1}^m \frac{\exp[\mathbf{q}_1\mathbf{k}_i^\top]}{\sum_j \exp[\mathbf{q}_1\mathbf{k}_j^\top]}\mathbf{v}_i} & \text{row vector 1 in } \mathbb{R}^{d_v} \\ \vdots \\ \underbrace{\sum_{i=1}^m \frac{\exp[\mathbf{q}_n\mathbf{k}_i^\top]}{\sum_j \exp[\mathbf{q}_n\mathbf{k}_j^\top]}\mathbf{v}_i} & \text{row vector } n \text{ in } \mathbb{R}^{d_v} \end{bmatrix}}_{n \times d_v}
\end{aligned}
\tag{9}
$$

looking at the i-th row of $\mathbf{A}$, i.e., $\mathbf{a}_i$, it can be expressed as a weighted sum of the rows of $\mathbf{V}$.

$$
\mathbf{a}_i\mathbf{V} = \begin{bmatrix} a_{i,1} & \dots & a_{i,m} \end{bmatrix} \begin{bmatrix} - & \mathbf{v}_1 & - \\ \vdots & \ddots & \vdots \\ - & \mathbf{v}_m & - \end{bmatrix} = \sum_{i=1}^m a_{i,j}\mathbf{v}_i
\tag{10}
$$

The output $\mathbf{O}$'s dimension has number of row to be the same as $\mathbf{Q}$, and column dimension to be d_V. This means that it has the number of elements of $\mathbf{Q}$, and each element is the same dimension as $\mathbf{v}_i$.

3

2 Self-attention and Multi-head Attention

The Transformer architecture uses multi-head attention, which is a key component of the Transformer architecture.Generally, input vectors are $(\mathbf{Q}, \mathbf{K}, \mathbf{V})$

When we let:

$$\mathbf{Q} = \mathbf{X}\mathbf{W}_Q, \quad \mathbf{K} = \mathbf{X}\mathbf{W}_K, \quad \mathbf{V} = \mathbf{X}\mathbf{W}_V \tag{11}$$

and let:

$$m = n = T \tag{12}$$

we achieved self-attention, in summary:

1. for each head i, linear transform $\mathbf{X}$ into several lower dimensional spaces via projection matrices $\mathbf{W}_i^q, \mathbf{K}\mathbf{W}_i^k, \mathbf{V}\mathbf{W}_i^v$ to obtain:

$$\mathbf{O}^{(i)} = \text{SDPA}(\mathbf{X}\mathbf{W}_Q^{(i)}, \mathbf{X}\mathbf{W}_K^{(i)}, \mathbf{X}\mathbf{W}_V^{(i)}) \tag{13}$$

2. each iteration i correspond to one "layer" of the ["linear", "Scaled Dot-Product Attention"] on the figure 1

3. then concatenate to produce output:

   with multi-head attention, we have multiple $\mathbf{O}_i$s therefore, we need to mix them together:

$$\mathbf{Y} = \underbrace{\text{concat}(\mathbf{O}_1, \ldots, \mathbf{O}_h)}_{T \times d_{\text{model}}} \underbrace{\mathbf{W}_O}_{d_{\text{model}} \times d_{\text{model}}} \tag{14}$$

2.1 Detailed implementation

2.1.1 class CausalSelfAttention(nn.Module)

```
class CausalSelfAttention(nn.Module):
```

this class defines the causal self-attention mechanism, the smallest unit of the Transformer architecture. We know that if we have h heads, then when we put multiple heads together, we have:

$$
\mathbf{Q} = \underbrace{\begin{bmatrix} - & \mathbf{x}_1 & - \\ - & \ldots & - \\ - & \mathbf{x}_T & - \end{bmatrix}}_{T \times d_{\text{model}}} \underbrace{\begin{bmatrix} \underbrace{\mathbf{W}_Q^{(1)}}_{d_{\text{model}} \times d_k} & \cdots & \underbrace{\mathbf{W}_Q^{(h)}}_{d_{\text{model}} \times d_k} \end{bmatrix}}_{d_{\text{model}} \times d_{\text{model}}} = \underbrace{\begin{bmatrix} \underbrace{\mathbf{Q}^{(1)}}_{T \times d_k} & \cdots & \underbrace{\mathbf{Q}^{(h)}}_{T \times d_k} \end{bmatrix}}_{T \times d_{\text{model}}}
$$

$$
\mathbf{K} = \underbrace{\begin{bmatrix} - & \mathbf{x}_1 & - \\ - & \ldots & - \\ - & \mathbf{x}_T & - \end{bmatrix}}_{T \times d_{\text{model}}} \underbrace{\begin{bmatrix} \underbrace{\mathbf{W}_K^{(1)}}_{d_{\text{model}} \times d_k} & \cdots & \underbrace{\mathbf{W}_K^{(h)}}_{d_{\text{model}} \times d_k} \end{bmatrix}}_{d_{\text{model}} \times d_{\text{model}}} = \underbrace{\begin{bmatrix} \underbrace{\mathbf{K}^{(1)}}_{T \times d_k} & \cdots & \underbrace{\mathbf{K}^{(h)}}_{T \times d_k} \end{bmatrix}}_{T \times d_{\text{model}}}
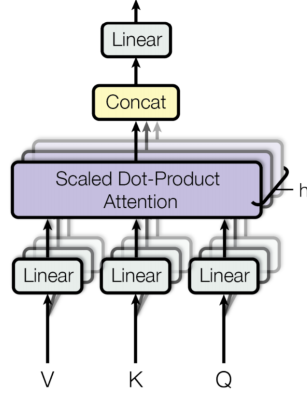\tag{15}
$$

4

Figure 1: Multi-head Attention

note that this also correpsonds to the B, T, C format (after adding the batch dimension B). where,

$$h \equiv n_{\text{heads}}$$
$$d_k = \text{d}_v$$
$$d_k \times h \equiv \text{d}_{\text{model}}$$
$$d_v \times h \equiv \text{d}_{\text{model}}$$

(16)

and the code has a few definitions:

```python
def __init__(self, d_model: int, n_heads: int, dropout: float, max_seq_len: int):
        super().__init__()
        assert d_model % n_heads == 0
        self.n_heads = n_heads
        self.head_dim = d_model // n_heads  # per-head dimensionality

        # Single projection that produces query, key, and value in one matmul
        self.qkv = nn.Linear(d_model, 3 * d_model)
        # Final projection after concatenating heads
        self.proj = nn.Linear(d_model, d_model)
        self.attn_dropout = nn.Dropout(dropout)
        self.resid_dropout = nn.Dropout(dropout)
```

Precompute variable mask matrix and store lower-triangular mask of shape $(1, 1, T, T)$, so it can be broadcasted over (B, h, T, T):

for example, for $T = 5$, we have:

$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 1 & 1 & 1 & 1 & 0 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

(17)

5

```
mask = torch.tril(torch.ones(max_seq_len, max_seq_len))
mask = mask.view(1, 1, max_seq_len, max_seq_len)
```

2.1.2 Step of each Transformer Layer

```
def forward(self, x: torch.Tensor) -> torch.Tensor:
```

1. At the conclusion of each Transformer layer, we have the output $\mathbf{Y}$ of shape (B, T, C), where it looks like (assume $B = 1$). See Eq.(29).

$$
\begin{bmatrix}
\underbrace{\begin{bmatrix}
- & \mathbf{y}_1 & - \\
& \vdots & \\
- & \mathbf{y}_T & -
\end{bmatrix}}_{T \times d_{\text{model}}}
\end{bmatrix}
\tag{18}
$$

The column dimesion is $h \times d_v$, which is also the size of the input vector $\mathbf{x}_t$. It can be thought as the "essential" form of the qkv matrix.

Now, we need to perform:

$$
\mathbf{Q} = \mathbf{X}\mathbf{W}_Q, \quad \mathbf{K} = \mathbf{X}\mathbf{W}_K, \quad \mathbf{V} = \mathbf{X}\mathbf{W}_V
\tag{19}
$$

to obtain the following single large matrix before splitting into $\mathbf{Q}, \mathbf{K}, \mathbf{V}$ matrices, we still assume $B = 1$.

$$
\begin{bmatrix}
\underbrace{\begin{bmatrix}
-\mathbf{q}_1^{(1)}- & \cdots & -\mathbf{q}_1^{(h)}- & -\mathbf{k}_1^{(1)}- & \cdots & -\mathbf{k}_1^{(h)}- & -\mathbf{v}_1^{(1)}- & \cdots & -\mathbf{v}_1^{(h)}- \\
\vdots & \ddots & \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\
-\mathbf{q}_T^{(1)}- & \cdots & -\mathbf{q}_T^{(h)}- & -\mathbf{k}_T^{(1)}- & \cdots & -\mathbf{k}_T^{(h)}- & -\mathbf{v}_T^{(1)}- & \cdots & -\mathbf{v}_T^{(h)}-
\end{bmatrix}}_{T \times (d_{\text{model}} \times 3)}
\end{bmatrix}
\tag{20}
$$

```
B, T, C = x.size()   # batch, time, channels
qkv = self.qkv(x)
```

after running the following code, we split the (B, T, C) format into (B, T, h, d_v), i.e., the big multi-head matrix containing all $\mathbf{Q}, \mathbf{K}, \mathbf{V}$ into separate $\mathbf{Q}, \mathbf{K}, \mathbf{V}$ matrices:

$$
\begin{bmatrix}
\underbrace{\begin{bmatrix}
- & \mathbf{q}_1^{(1)} & - & \cdots & - & \mathbf{q}_1^{(h)} & - \\
& \vdots & & \ddots & & \vdots & \\
- & \mathbf{q}_T^{(1)} & - & \cdots & - & \mathbf{q}_T^{(h)} & -
\end{bmatrix}}_{T \times d_{\text{model}}}
\underbrace{\begin{bmatrix}
- & \mathbf{k}_1^{(1)} & - & \cdots & - & \mathbf{k}_1^{(h)} & - \\
& \vdots & & \ddots & & \vdots & \\
- & \mathbf{k}_T^{(1)} & - & \cdots & - & \mathbf{k}_T^{(h)} & -
\end{bmatrix}}_{T \times d_{\text{model}}}
\underbrace{\begin{bmatrix}
- & \mathbf{v}_1^{(1)} & - & \cdots & - & \mathbf{v}_1^{(h)} & - \\
& \vdots & & \ddots & & \vdots & \\
- & \mathbf{v}_T^{(1)} & - & \cdots & - & \mathbf{v}_T^{(h)} & -
\end{bmatrix}}_{T \times d_{\text{model}}}
\end{bmatrix}
\tag{21}
$$

```
        q, k, v = qkv.split(C, dim=2)  # each of shape (B, T, C)
```

2. now we split the **QKV** single large matrix into each separate $\mathbf{Q}, \mathbf{K}, \mathbf{V}$ matrices:

```
def reshape_heads(t: torch.Tensor) -> torch.Tensor:
        # Reshape to (B, n_heads, T, head_dim) for parallel head
            computation
        return t.view(B, T, self.n_heads, self.head_dim).transpose(1, 2)

# (B, n_heads, T, head_dim)
q = reshape_heads(q)
k = reshape_heads(k)
v = reshape_heads(v)
```

after running the above code, it has split up $\mathbf{Q}$, $\mathbf{K}$ and $\mathbf{V}$, so that each will be in the form of for simplicity let $B = 1$, and we would like to have the structure to look like:

$$
\mathbf{Q} = \underbrace{\left[\underbrace{\begin{bmatrix} - & \mathbf{q}_1 & - \\ - & \dots & - \\ - & \mathbf{q}_T & - \end{bmatrix}}_{\mathbf{Q}_{(1)}:\, T \times d_k} \quad \dots \quad \underbrace{\begin{bmatrix} - & \mathbf{q}_1 & - \\ - & \dots & - \\ - & \mathbf{q}_T & - \end{bmatrix}}_{\mathbf{Q}_{(h)}:\, T \times d_k} \right]}_{T \times d_{\mathrm{model}}} \quad
\mathbf{K} = \underbrace{\left[\underbrace{\begin{bmatrix} - & \mathbf{k}_1 & - \\ - & \dots & - \\ - & \mathbf{k}_T & - \end{bmatrix}}_{\mathbf{K}_{(1)}:\, T \times d_k} \quad \dots \quad \underbrace{\begin{bmatrix} - & \mathbf{k}_1 & - \\ - & \dots & - \\ - & \mathbf{k}_T & - \end{bmatrix}}_{\mathbf{K}_{(h)}:\, T \times d_k} \right]}_{T \times d_{\mathrm{model}}}
$$

$$
\mathbf{V} = \underbrace{\left[\underbrace{\begin{bmatrix} - & \mathbf{v}_1 & - \\ - & \dots & - \\ - & \mathbf{v}_T & - \end{bmatrix}}_{\mathbf{V}_{(1)}:\, T \times d_v} \quad \dots \quad \underbrace{\begin{bmatrix} - & \mathbf{v}_1 & - \\ - & \dots & - \\ - & \mathbf{v}_T & - \end{bmatrix}}_{\mathbf{V}_{(h)}:\, T \times d_v} \right]}_{T \times d_{\mathrm{model}}}
\tag{22}
$$

where:

$$
\begin{aligned}
h &\equiv n_{\mathrm{heads}} \\
d_k \times h &\equiv \mathrm{d}_{\mathrm{model}} \\
d_v \times h &\equiv \mathrm{d}_{\mathrm{model}}
\end{aligned}
\tag{23}
$$

3. we then apply $\mathbf{Q}\mathbf{K}^\top$:

```
att = (q @ k.transpose(-2, -1)) # (B, n_heads, T, T)
```

note the above code does not perform $\mathbf{Q}\mathbf{K}^\top$, after $\mathbf{Q}$ and $\mathbf{K}$ are concatenated from all the heads, but rather it performs $\mathbf{Q}\mathbf{K}^\top$ for each head, i.e., broadcasted matrix multiplication:

$$
\begin{bmatrix} \mathbf{Q}_{(1)}\mathbf{K}_{(1)}^\top & \dots & \mathbf{Q}_{(h)}\mathbf{K}_{(h)}^\top \end{bmatrix}
\tag{24}
$$

7

4. we then scale by $1/\sqrt{d_k}$:

```
att = att / (self.head_dim ** 0.5)
```

where we obtained:

$$
\left[\frac{\mathbf{Q}_{(1)}\mathbf{K}_{(1)}^{\top}}{\sqrt{d_k}} \quad \cdots \quad \frac{\mathbf{Q}_{(h)}\mathbf{K}_{(h)}^{\top}}{\sqrt{d_k}}\right] = \left[\underbrace{\begin{bmatrix} \frac{\mathbf{q}_1^{(1)}\mathbf{k}_1^{(1)\top}}{\sqrt{d_k}} & \cdots & \frac{\mathbf{q}_1^{(1)}\mathbf{k}_T^{(1)\top}}{\sqrt{d_k}} \\ \vdots & \ddots & \vdots \\ \frac{\mathbf{q}_T^{(1)}\mathbf{k}_1^{(1)\top}}{\sqrt{d_k}} & \cdots & \frac{\mathbf{q}_T^{(1)}\mathbf{k}_T^{(1)\top}}{\sqrt{d_k}} \end{bmatrix}}_{\frac{\mathbf{Q}^{(1)}\mathbf{K}^{(1)\top}}{\sqrt{d_k}}} \quad \cdots \quad \underbrace{\begin{bmatrix} \frac{\mathbf{q}_1^{(h)}\mathbf{k}_1^{(h)\top}}{\sqrt{d_k}} & \cdots & \frac{\mathbf{q}_1^{(h)}\mathbf{k}_T^{(h)\top}}{\sqrt{d_k}} \\ \vdots & \ddots & \vdots \\ \frac{\mathbf{q}_T^{(h)}\mathbf{k}_1^{(h)\top}}{\sqrt{d_k}} & \cdots & \frac{\mathbf{q}_T^{(h)}\mathbf{k}_T^{(h)\top}}{\sqrt{d_k}} \end{bmatrix}}_{\mathbf{Q}^{(h)}\mathbf{K}^{(h)\top}}\right]
\tag{25}
$$

5. we apply a causal mask to the attention matrix, so that each word can only attend to the previous words:

```
# Apply causal mask so position t cannot attend to positions > t
att = att.masked_fill(self.mask[:, :, :T, :T] == 0, float('-inf'))
```

$$
\left[\begin{bmatrix} \frac{\mathbf{q}_1^{(1)}\mathbf{k}_1^{(1)\top}}{\sqrt{d_k}} & \cdots & -\infty \\ \vdots & \ddots & \vdots \\ \frac{\mathbf{q}_T^{(1)}\mathbf{k}_1^{(1)\top}}{\sqrt{d_k}} & \cdots & \frac{\mathbf{q}_T^{(1)}\mathbf{k}_T^{(1)\top}}{\sqrt{d_k}} \end{bmatrix} \quad \cdots \quad \begin{bmatrix} \frac{\mathbf{q}_1^{(h)}\mathbf{k}_1^{(h)\top}}{\sqrt{d_k}} & \cdots & -\infty \\ \vdots & \ddots & \vdots \\ \frac{\mathbf{q}_T^{(h)}\mathbf{k}_1^{(h)\top}}{\sqrt{d_k}} & \cdots & \frac{\mathbf{q}_T^{(h)}\mathbf{k}_T^{(h)\top}}{\sqrt{d_k}} \end{bmatrix}\right]
\tag{26}
$$

6. apply row-wise softmax to the attention matrix:

```
att = F.softmax(att, dim=-1)
att = self.attn_dropout(att)
```

$$
\left[\underbrace{\begin{bmatrix} \mathrm{softmax}\left(\begin{bmatrix} \frac{\mathbf{q}_1^{(1)}\mathbf{k}_1^{(1)\top}}{\sqrt{d_k}} & \cdots & -\infty \end{bmatrix}\right) \\ \vdots \\ \mathrm{softmax}\left(\begin{bmatrix} \frac{\mathbf{q}_T^{(1)}\mathbf{k}_1^{(1)\top}}{\sqrt{d_k}} & \cdots & \frac{\mathbf{q}_T^{(1)}\mathbf{k}_T^{(1)\top}}{\sqrt{d_k}} \end{bmatrix}\right) \end{bmatrix}}_{\mathrm{softmax}\left(\frac{\mathbf{Q}^{(1)}\mathbf{K}^{(1)\top}}{\sqrt{d_k}}\right)} \quad \cdots \quad \underbrace{\begin{bmatrix} \mathrm{softmax}\left(\begin{bmatrix} \frac{\mathbf{q}_1^{(h)}\mathbf{k}_1^{(h)\top}}{\sqrt{d_k}} & \cdots & -\infty \end{bmatrix}\right) \\ \vdots \\ \mathrm{softmax}\left(\begin{bmatrix} \frac{\mathbf{q}_T^{(h)}\mathbf{k}_1^{(h)\top}}{\sqrt{d_k}} & \cdots & \frac{\mathbf{q}_T^{(h)}\mathbf{k}_T^{(h)\top}}{\sqrt{d_k}} \end{bmatrix}\right) \end{bmatrix}}_{\mathrm{softmax}\left(\frac{\mathbf{Q}^{(h)}\mathbf{K}^{(h)\top}}{\sqrt{d_k}}\right)}\right]
\tag{27}
$$

7. apply weighted sum to the values:

```
# Weighted sum of values
y = att @ v   # (B, n_heads, T, head_dim)
```

where we obtained:

$$\underbrace{\left[\underbrace{\mathrm{softmax}\left(\frac{\mathbf{Q}^{(1)}\mathbf{K}^{(1)\top}}{\sqrt{d_k}}\right)\mathbf{V}^{(1)}}_{T\times d_v} \quad \dots \quad \underbrace{\mathrm{softmax}\left(\frac{\mathbf{Q}^{(h)}\mathbf{K}^{(h)\top}}{\sqrt{d_k}}\right)\mathbf{V}^{(h)}}_{T\times d_v}\right]}_{T\times d_{\mathrm{model}}} \tag{28}$$

8. transform the output back to the original dimension B, T, C, where contiguous() is a method that returns a contiguous tensor.

```
# (B, n_heads, T, head_dim) -> (B, T, n_heads, head_dim)
y = y.transpose(1, 2)
y = y.contiguous()
# Reassemble heads: (B, T, C)
y = y.view(B, T, C)
```

know that:

$$
\underbrace{\begin{bmatrix} - & \mathbf{y}_1 & - \\ & \vdots & \\ - & \mathbf{y}_T & - \end{bmatrix}}_{T\times d_{\mathrm{model}}} = \underbrace{\left[\underbrace{\begin{bmatrix} - & \mathbf{o}_1^{(1)} & - \\ - & \dots & - \\ - & \mathbf{o}_T^{(1)} & - \end{bmatrix}}_{\mathbf{O}_{(1)}:\, T\times d_v} \quad \dots \quad \underbrace{\begin{bmatrix} - & \mathbf{o}_1^{(h)} & - \\ - & \dots & - \\ - & \mathbf{o}_T^{(h)} & - \end{bmatrix}}_{\mathbf{O}_{(h)}:\, T\times d_v}\right]}_{T\times d_{\mathrm{model}}} \underbrace{\mathbf{W}_O}_{d_{\mathrm{model}}\times d_{\mathrm{model}}}
$$

$$
= \underbrace{\begin{bmatrix} - & \mathbf{o}_1^{(1)} & - & \dots & - & \mathbf{o}_1^{(h)} & - \\ & \vdots & & \ddots & & \vdots & \\ - & \mathbf{o}_T^{(1)} & - & \dots & - & \mathbf{o}_T^{(h)} & - \end{bmatrix}}_{T\times d_{\mathrm{model}}} \underbrace{\mathbf{W}_O}_{d_{\mathrm{model}}\times d_{\mathrm{model}}} \tag{29}
$$

each row of $\mathbf{y}_i$ can be determined independently and in parallel.

9. we then perform a final linear projection on the concatenated heads:

```
# Final linear projection on the concatenated heads
y = self.resid_dropout(self.proj(y)) # resid_dropout is the residual
    dropout
return y
```

one can see that the final output is a tensor of shape B, T, d_{model}.

2.1.3   class TransformerBlock(nn.Module)

The following is the code for Pre-norm Transformer block with attention followed by an MLP, where the CausalSelfAttention() is from the previous section:

```python
class TransformerBlock(nn.Module):
    def __init__(self, d_model: int, n_heads: int, d_ff: int, dropout: float,
        max_seq_len: int):
        super().__init__()
        # Pre-layernorm improves optimization stability in autoregressive LMs
        self.ln1 = nn.LayerNorm(d_model)
        self.attn = CausalSelfAttention(d_model, n_heads, dropout, max_seq_len)
        self.ln2 = nn.LayerNorm(d_model)
        # Feed-forward network (position-wise MLP)
        self.mlp = nn.Sequential(
            nn.Linear(d_model, d_ff),
            nn.GELU(),
            nn.Dropout(dropout),
            nn.Linear(d_ff, d_model),
            nn.Dropout(dropout),
        )
    def forward(self, x: torch.Tensor) -> torch.Tensor:
        # Attention sub-layer with residual connection
        x = x + self.attn(self.ln1(x))
        # MLP sub-layer with residual connection
        x = x + self.mlp(self.ln2(x))
        return x
```

$x = x + f(x)$ is the residual connection. and nn.GELU() is the activation function.

2.1.4 class TransformerLM(nn.Module)

```python
class TransformerLM(nn.Module):
    """Causal language model built from Transformer blocks."""

    def __init__(self, config: ModelConfig):
        super().__init__()
        self.config = config
        # Token and positional embeddings are summed to form the input stream
        self.token_emb = nn.Embedding(config.vocab_size, config.d_model)
        self.pos_emb = nn.Embedding(config.max_seq_len, config.d_model)
        self.drop = nn.Dropout(config.dropout)
        # Stack of identical Transformer blocks
        self.blocks = nn.ModuleList([
            TransformerBlock(config.d_model, config.n_heads, config.d_ff, config
                .dropout, config.max_seq_len)
            for _ in range(config.n_layers)
        ])
        # Final layernorm before projecting to vocabulary logits
        self.ln_f = nn.LayerNorm(config.d_model)
        # LM head ties to embedding size but does not share weights here
        self.lm_head = nn.Linear(config.d_model, config.vocab_size, bias=False)

    def forward(self, idx: torch.Tensor, targets: Optional[torch.Tensor] = None)
        :
```

```
"""Forward pass producing next-token logits and optional cross-entropy
    loss.

Parameters
----------
idx : torch.Tensor
    Token indices of shape (B, T).
targets : Optional[torch.Tensor]
    Target indices of shape (B, T). If provided, a token-level
    cross-entropy loss is computed and returned.

Returns
-------
Tuple[torch.Tensor, Optional[torch.Tensor]]
    - logits: (B, T, vocab_size)
    - loss: scalar loss if targets is provided, else None
"""
B, T = idx.shape
# Ensure sequence length fits within the configured context window
assert T <= self.config.max_seq_len

# Positions 0..T-1 for current sequence length
pos = torch.arange(0, T, dtype=torch.long, device=idx.device).unsqueeze
    (0)
```

this is to create a tensor of shape $(1, T)$ with values from 0 to T-1. We call the unsqueeze(0) to add a batch dimension, so that it can be broadcasted over the batch dimension when adding to the token embeddings.

$$\begin{bmatrix} \begin{bmatrix} 0 & 1 & 2 & \ldots & T-1 \end{bmatrix} \end{bmatrix} \tag{30}$$

```
# Embed tokens and positions, then apply dropout
x = self.token_emb(idx) + self.pos_emb(pos)
x = self.drop(x)

# Pass through the stack of Transformer blocks
for blk in self.blocks:
    x = blk(x)

# Final normalization and projection to vocabulary
x = self.ln_f(x)
logits = self.lm_head(x)

# Compute token-level cross-entropy if targets are given
loss = None
if targets is not None:
    loss = F.cross_entropy(logits.view(-1, logits.size(-1)), targets.
        view(-1))
return logits, loss
```

```python
@torch.no_grad()
def generate(self, idx: torch.Tensor, max_new_tokens: int) -> torch.Tensor:
    """Greedy-free sampling generation from the language model.

    Iteratively samples 'max_new_tokens' next tokens by:
    - conditioning on the last 'max_seq_len' tokens
    - computing logits for the next position
    - sampling from the softmax distribution
    - appending the sampled token to the context

    Parameters
    ----------
    idx : torch.Tensor
        Initial context token indices of shape (B, T0).
    max_new_tokens : int
        Number of new tokens to sample and append.

    Returns
    -------
    torch.Tensor
        Extended token indices of shape (B, T0 + max_new_tokens).
    """
    self.eval()
    for _ in range(max_new_tokens):
        # Restrict conditioning length to model's maximum context window
        # only the last self.config.max_seq_len tokens are used to condition
        #     the generation, as they are the most recent tokens
        idx_cond = idx[:, -self.config.max_seq_len:]
        logits, _ = self(idx_cond)
        # Select logits for the last time step (the next token prediction)
        logits = logits[:, -1, :]
        probs = F.softmax(logits, dim=-1)
        # Multinomial sampling allows for non-greedy, stochastic generation
        next_id = torch.multinomial(probs, num_samples=1)
        # Append sampled token and continue
        idx = torch.cat((idx, next_id), dim=1)
    return idx
```

2.2 Cross-Attention

If we were to perform cross-attention, $\mathbf{K}$ and $\mathbf{V}$ must be the same thing. For example,

1. Text-to-image Cross-Attention:

$$
\begin{aligned}
&\mathbf{K} \text{ and } \mathbf{V} \text{ are text embedding} \\
&\mathbf{Q} \text{ is image embedding}
\end{aligned}
\tag{31}
$$

2. Image-to-text Cross-Attention:

$$\mathbf{K} \text{ and } \mathbf{V} \text{ are image embedding}$$
$$\mathbf{Q} \text{ is text embedding} \tag{32}$$

3 $\mathbf{K} - \mathbf{V}$ caching

looking at $\mathbf{QK}^\top$, imagine in a self-attention scenario, where the current senetence length is T, and $m = n = T$, then $\mathbf{QK}^\top$ is a square matrix of $T \times T$, where we have:

$$\mathbf{QK}^\top = \begin{bmatrix} \mathbf{q}_1\mathbf{k}_1^\top & \dots & \mathbf{q}_1\mathbf{k}_T^\top \\ \vdots & \ddots & \vdots \\ \mathbf{q}_T\mathbf{k}_1^\top & \dots & \mathbf{q}_T\mathbf{k}_T^\top \end{bmatrix} \tag{33}$$

let's say we have a new word w_{T+1} to be added to the sentence, then we have:

$$\begin{bmatrix} \mathbf{q}_1\mathbf{k}_1^\top & \dots & \mathbf{q}_1\mathbf{k}_T^\top & \mathbf{q}_1\mathbf{k}_{T+1}^\top \\ \vdots & \ddots & \vdots & \vdots \\ \mathbf{q}_T\mathbf{k}_1^\top & \dots & \mathbf{q}_T\mathbf{k}_T^\top & \mathbf{q}_T\mathbf{k}_{T+1}^\top \\ \mathbf{q}_{T+1}\mathbf{k}_1^\top & \dots & \mathbf{q}_{T+1}\mathbf{k}_T^\top & \mathbf{q}_{T+1}\mathbf{k}_{T+1}^\top \end{bmatrix} = \begin{bmatrix} \mathbf{QK}^\top & \mathbf{Qk}_{T+1}^\top \\ \mathbf{q}_{T+1}\mathbf{K}^\top & \mathbf{q}_{T+1}\mathbf{k}_{T+1}^\top \end{bmatrix} \tag{34}$$

From here, one sees that in generic matrix multiplication, one needs to cache both $\mathbf{K}$ and $\mathbf{Q}$. However, the matrix masks out the upper triangular part of the matrix, i.e., it looks like this:

$$\begin{bmatrix} \mathbf{q}_1\mathbf{k}_1^\top & - & - & - \\ \vdots & \ddots & \vdots & \vdots \\ \mathbf{q}_T\mathbf{k}_1^\top & \dots & \mathbf{q}_T\mathbf{k}_T^\top & - \\ \mathbf{q}_{T+1}\mathbf{k}_1^\top & \dots & \mathbf{q}_{T+1}\mathbf{k}_T^\top & \mathbf{q}_{T+1}\mathbf{k}_{T+1}^\top \end{bmatrix} = \begin{bmatrix} \underbrace{\mathbf{QK}^\top}_{\text{lower triangular}} & - \\ \mathbf{q}_{T+1}\mathbf{K}^\top & \mathbf{q}_{T+1}\mathbf{k}_{T+1}^\top \end{bmatrix} \tag{35}$$

One can see the newly added row (no new column is added) in the above matrix, requires $\mathbf{q}_{T+1}$, $\mathbf{k}_{T+1}$ and $\mathbf{K}$. This means that $\mathbf{K}$ needs to be cached.

Again looking at Eq.(9):

$$\begin{bmatrix} \frac{\exp(\mathbf{q}_1\mathbf{k}_1^\top)}{\exp(\mathbf{q}_1\mathbf{k}_1^\top)}\mathbf{v}_1 \\ \\ \frac{\exp(\mathbf{q}_2\mathbf{k}_1^\top)}{\exp(\mathbf{q}_2\mathbf{k}_1^\top)+\exp(\mathbf{q}_2\mathbf{k}_2^\top)}\mathbf{v}_1 + \frac{\exp(\mathbf{q}_2\mathbf{k}_2^\top)}{\exp(\mathbf{q}_2\mathbf{k}_1^\top)+\exp(\mathbf{q}_2\mathbf{k}_2^\top)}\mathbf{v}_2 \\ \\ \frac{\exp(\mathbf{q}_3\mathbf{k}_1^\top)}{\exp(\mathbf{q}_3\mathbf{k}_1^\top)+\exp(\mathbf{q}_3\mathbf{k}_2^\top)+\exp(\mathbf{q}_3\mathbf{k}_3^\top)}\mathbf{v}_1 + \frac{\exp(\mathbf{q}_3\mathbf{k}_2^\top)}{\exp(\mathbf{q}_3\mathbf{k}_1^\top)+\exp(\mathbf{q}_3\mathbf{k}_2^\top)+\exp(\mathbf{q}_3\mathbf{k}_3^\top)}\mathbf{v}_2 + \frac{\exp(\mathbf{q}_3\mathbf{k}_3^\top)}{\exp(\mathbf{q}_3\mathbf{k}_1^\top)+\exp(\mathbf{q}_3\mathbf{k}_2^\top)+\exp(\mathbf{q}_3\mathbf{k}_3^\top)}\mathbf{v}_3 \\ \vdots \\ \sum_{i=1}^{T} \frac{\exp(\mathbf{q}_T\mathbf{k}_i^\top)}{\sum_{j=1}^{T}\exp(\mathbf{q}_T\mathbf{k}_j^\top)}\mathbf{v}_i \\ \\ \sum_{i=1}^{T+1} \frac{\exp(\mathbf{q}_{T+1}\mathbf{k}_i^\top)}{\sum_{j=1}^{T+1}\exp(\mathbf{q}_{T+1}\mathbf{k}_j^\top)}\mathbf{v}_i \end{bmatrix} \tag{36}$$

the last row in the above matrix contains the newly added row, where all $\mathbf{v}_i$ are needed for computation and hence they need to be cached.

3.0.1 Multi-head latent attention

say we have sentence of length T, and we have each word represented by a $d = 7168$ dimensional vector. Then we can have:

$$
\begin{aligned}
\mathbf{L}_{K,V} &= \mathbf{X}\mathbf{W}_{\downarrow K,V} \\
&= \underbrace{\begin{bmatrix} - & \mathbf{x}_1 & - \\ - & \dots & - \\ - & \mathbf{x}_T & - \end{bmatrix}}_{T \times d_{\text{model}}} \underbrace{\begin{bmatrix} \mathbf{w}_{1,1} & \cdots & \mathbf{w}_{1,576} \\ \vdots & \ddots & \vdots \\ \mathbf{w}_{7168,1} & \cdots & \mathbf{w}_{7168,576} \end{bmatrix}}_{d_{\text{model}} \times 576}
\end{aligned}
\tag{37}
$$

where $\mathbf{L}_{K,V}$ is the "common" part of the latent representation of the sentence for both $\mathbf{K}$ and $\mathbf{V}$.

normally, in the case of self-attention, one should have, $\mathbf{Q} = \mathbf{X}\mathbf{W}_Q$, $\mathbf{K} = \mathbf{X}\mathbf{W}_K$, $\mathbf{V} = \mathbf{X}\mathbf{W}_V$, but in here, we are having:

$$
\begin{array}{lcl}
\text{traditional} & & \text{multi-head latent attention} \\[2mm]
\mathbf{Q} = \mathbf{X}\mathbf{W}_Q & \rightarrow & \underbrace{\mathbf{Q}}_{T \times 128} = \mathbf{X} \underbrace{\mathbf{W}_Q}_{7168 \times 128} \\[4mm]
\mathbf{K} = \mathbf{X}\mathbf{W}_K & \rightarrow & \underbrace{\mathbf{K}}_{T \times 128} = \mathbf{X} \underbrace{\mathbf{W}_{\downarrow K,V}}_{d_{\text{model}} \times 576} \underbrace{\mathbf{W}_{\uparrow K}}_{576 \times 128} = \underbrace{\mathbf{L}_{K,V}}_{T \times 576} \underbrace{\mathbf{W}_{\uparrow K}}_{576 \times 128} \\[4mm]
\mathbf{V} = \mathbf{X}\mathbf{W}_V & \rightarrow & \underbrace{\mathbf{V}}_{T \times 128} = \mathbf{X} \underbrace{\mathbf{W}_{\downarrow K,V}}_{d_{\text{model}} \times 576} \underbrace{\mathbf{W}_{\uparrow V}}_{576 \times 128} = \underbrace{\mathbf{L}_{K,V}}_{T \times 576} \underbrace{\mathbf{W}_{\uparrow V}}_{576 \times 128}
\end{array}
\tag{38}
$$

making $\mathbf{W}_{\downarrow K,V}$ to be common for all heads for both $\mathbf{K}$ and $\mathbf{V}$. Therefore only a smaller $\mathbf{W}_{\uparrow K}$ and $\mathbf{W}_{\uparrow V}$ need to be trained per head for both $\mathbf{K}$ and $\mathbf{V}$. This means that:

$$
\begin{aligned}
\mathbf{W}_K &\longrightarrow \mathbf{W}_{\downarrow K,V} \, \mathbf{W}_{\uparrow K} \\
\mathbf{K} = \mathbf{W}_K &\longrightarrow \mathbf{X}\mathbf{W}_{\downarrow K,V} \, \mathbf{W}_{\uparrow K}
\end{aligned}
\tag{39}
$$

substitute:

$$
\begin{aligned}
\mathbf{Q}\mathbf{K}^\top &= \mathbf{X}\mathbf{W}_Q(\mathbf{X} \, \mathbf{W}_{\downarrow K,V} \, \mathbf{W}_{\uparrow K})^\top \\
&= \mathbf{X}\mathbf{W}_Q\big((\mathbf{X} \, \mathbf{W}_{\downarrow K,V}) \, \mathbf{W}_{\uparrow K})\big)^\top \\
&= \underbrace{\mathbf{X}}_{T \times d_{\text{model}}} \underbrace{(\mathbf{W}_Q\mathbf{W}_{\uparrow K}^\top)}_{d_{\text{model}} \times 576} \underbrace{(\mathbf{X} \, \mathbf{W}_{\downarrow K,V})^\top}_{576 \times T}
\end{aligned}
\tag{40}
$$

14

$$
\begin{aligned}
\mathbf{Q}\mathbf{K}^\top &= \mathbf{X}\mathbf{W}_Q(\mathbf{X}\ \mathbf{W}_{\downarrow K,V}\ \mathbf{W}_{\uparrow K})^\top \\
&= \mathbf{X}\mathbf{W}_Q\big((\mathbf{X}\ \mathbf{W}_{\downarrow K,V})\ \mathbf{W}_{\uparrow K})\big)^\top \\
&= \underbrace{\mathbf{X}}_{T\times d_{\text{model}}}\ \underbrace{(\mathbf{W}_Q\mathbf{W}_{\uparrow K}^\top)}_{d_{\text{model}}\times 576}\underbrace{(\mathbf{X}\ \mathbf{W}_{\downarrow K,V})^\top}_{576\times T}
\end{aligned}
\tag{41}
$$

- Remember that although we have different $\mathbf{W}_{\uparrow K}$ for each head, $(\mathbf{W}_Q\mathbf{W}_{\uparrow K}^\top)$ can be computed first only once, before they multiple with $\mathbf{X}$ (and $(\mathbf{X}\ \mathbf{W}_{\downarrow K,V})$) during inference time for each variable length $\mathbf{X}$. This means that we do not need to store $\mathbf{W}_Q$ and $\mathbf{W}_{\uparrow K}$ separately, but rather we can store $\mathbf{W}_Q\mathbf{W}_{\uparrow K}^\top$ once. It's like having a new $\mathbf{W}_Q$

In terms of caching, assumed that at the lenth of T words, we have cached $\mathbf{X}(\mathbf{W}_Q\mathbf{W}_{\uparrow K}^\top)$ and $\mathbf{X}\mathbf{W}_{\downarrow K,V}$. Then after adding a new word, i.e., $\mathbf{X}_{T+1}$, we can just compute one more row of the matrix $\mathbf{X}_{T+1}\mathbf{W}_Q\mathbf{W}_{\uparrow K}$ and $\mathbf{X}_{T+1}\mathbf{W}_{\downarrow K,V}$ and append it to the cache, i.e.,

$$
\begin{aligned}
&\begin{bmatrix} - & \mathbf{x}_1 & - \\ - & \dots & - \\ - & \mathbf{x}_T & - \end{bmatrix}\big[(\mathbf{W}_Q\mathbf{W}_{\uparrow K}^\top)\big] = \text{cache} \\
\Longrightarrow &\begin{bmatrix} - & \mathbf{x}_1 & - \\ - & \dots & - \\ - & \mathbf{x}_T & - \\ - & \mathbf{x}_{T+1} & - \end{bmatrix}\big[(\mathbf{W}_Q\mathbf{W}_{\uparrow K}^\top)\big] = \begin{bmatrix} \text{cache} \\ \mathbf{x}_{T+1}(\mathbf{W}_Q\mathbf{W}_{\uparrow K}^\top) \end{bmatrix}
\end{aligned}
\tag{42}
$$

the same applies to $\mathbf{X}\mathbf{W}_{\downarrow K,V}$, i.e.,

$$
\begin{aligned}
&\begin{bmatrix} - & \mathbf{x}_1 & - \\ - & \dots & - \\ - & \mathbf{x}_T & - \end{bmatrix}\big[\mathbf{W}_{\downarrow K,V}\big] = \text{cache} \\
\Longrightarrow &\begin{bmatrix} - & \mathbf{x}_1 & - \\ - & \dots & - \\ - & \mathbf{x}_T & - \\ - & \mathbf{x}_{T+1} & - \end{bmatrix}\big[\mathbf{W}_{\downarrow K,V}\big] = \begin{bmatrix} \text{cache} \\ \mathbf{x}_{T+1}\mathbf{W}_{\downarrow K,V} \end{bmatrix}
\end{aligned}
\tag{43}
$$

Now, let's look at the output of the multi-head attention:

$$\mathbf{Y} = \underbrace{\begin{bmatrix} - & \mathbf{y}_1 & - \\ & \vdots & \\ - & \mathbf{y}_T & - \end{bmatrix}}_{T \times d_{\text{model}}} = \underbrace{\text{concat}(\mathbf{O}_1, \ldots, \mathbf{O}_h)}_{T \times (d_v \times h)} \underbrace{\mathbf{W}_O}_{d_{\text{model}} \times d_{\text{model}}}$$

$$= \left[\text{softmax}\left(\frac{\mathbf{Q}_{(1)}\mathbf{K}_{(1)}^\top}{\sqrt{d_k}}\right)\mathbf{V}_{(1)} \quad \ldots \quad \text{softmax}\left(\frac{\mathbf{Q}_{(h)}\mathbf{K}_{(h)}^\top}{\sqrt{d_k}}\right)\mathbf{V}_{(h)} \right] \mathbf{W}_O$$

$$= \left[\mathbf{O}_{(1)}\mathbf{V}_{(1)} \quad \ldots \quad \mathbf{O}_{(h)}\mathbf{V}_{(h)} \right] \mathbf{W}_O$$

$$= \left[\underbrace{\mathbf{O}_{(1)}}_{T \times d_v} \quad \cdots \quad \underbrace{\mathbf{O}_{(h)}}_{T \times d_v} \right] \begin{bmatrix} \underbrace{\mathbf{V}_{(1)}}_{d_v \times d_v} & \cdots & \mathbf{0} \\ \vdots & \ddots & \vdots \\ \mathbf{0} & \cdots & \underbrace{\mathbf{V}_{(h)}}_{d_v \times d_v} \end{bmatrix} \begin{bmatrix} \underbrace{\mathbf{W}_O^{1:d_v, 1:d_{\text{model}}}}_{d_v \times d_{\text{model}}} \\ \vdots \\ \underbrace{\mathbf{W}_O^{d_v \times (h-1)+1:d_{\text{model}}, 1:d_{\text{model}}}}_{d_v \times d_{\text{model}}} \end{bmatrix} \quad (44)$$

$$= \underbrace{\left[\mathbf{O}_{(1)} \quad \ldots \quad \mathbf{O}_{(h)} \right]}_{T \times d_{\text{model}}} \underbrace{\begin{bmatrix} \underbrace{\mathbf{V}_{(1)}\mathbf{W}_O^{1:d_v, 1:d_{\text{model}}}}_{d_v \times d_{\text{model}}} \\ \vdots \\ \underbrace{\mathbf{V}_{(h)}\mathbf{W}_O^{d_v \times (h-1)+1:d_{\text{model}}, 1:d_{\text{model}}}}_{d_v \times d_{\text{model}}} \end{bmatrix}}_{d_{\text{model}} \times d_{\text{model}}}$$

now that we replaced $\mathbf{W}_V$ with $\mathbf{W}_{\downarrow K,V}\mathbf{W}_{\uparrow V}$, we can have:

$$\mathbf{V} = \underbrace{\mathbf{X}}_{T \times d_{\text{model}}} \underbrace{\mathbf{W}_{\downarrow K,V}}_{d_{\text{model}} \times 576} \underbrace{\mathbf{W}_{\uparrow V}}_{576 \times d_v} \quad (45)$$

therefore, we can have each $\mathbf{V}$, for example $\mathbf{V}_{(1)}$ as

$$\mathbf{V}_{(1)}\mathbf{W}_O^{1:d_v, 1:d_{\text{model}}} = \underbrace{\mathbf{X}}_{T \times d_{\text{model}}} \underbrace{\mathbf{W}_{\downarrow K,V}}_{d_{\text{model}} \times 576} \underbrace{\mathbf{W}_{\uparrow V}}_{576 \times d_v} \underbrace{\mathbf{W}_O^{1:d_v, 1:d_{\text{model}}}}_{d_v \times d_{\text{model}}} \quad (46)$$

we already computed $\mathbf{X}\mathbf{W}_{\downarrow K,V}$ as $\mathbf{L}_{K,V}$, and we can also pre-compute $\mathbf{W}_{\uparrow V}\mathbf{W}_O^{1:d_v, 1:d_{\text{model}}}$ as single matrix, instead of storing $\mathbf{W}_{\uparrow V}$ and $\mathbf{W}_O^{1:d_v, 1:d_{\text{model}}}$ separately.

Finally, we have $\mathbf{Y}$ as the output of the multi-head attention layer.

4 Rotary Embedding RoPE

Here is the RoPE formula:

```
class RotaryEmbedding(nn.Module):
    def __init__(self, head_dim=64, rope_theta=100000):
```

```
        super().__init__()
        self.head_dim = head_dim
        self.rope_theta = rope_theta

    def forward(self, x):
        # x shape: [batch, num_heads, seq_len, head_dim]
```

```
        # Compute frequencies
        i = torch.arange(0, self.head_dim, 2)
        freqs = 1.0 / (self.rope_theta ** (i / self.head_dim))
```

$$i = \begin{bmatrix} 0 & 2 & 4 & \ldots & d_v - 2 \end{bmatrix}$$
$$\text{freqs} = \underbrace{\begin{bmatrix} \frac{1}{\theta^{0/d_v}} & \frac{1}{\theta^{2/d_v}} & \frac{1}{\theta^{4/d_v}} & \cdots & \frac{1}{\theta^{(d_v-2)/d_v}} \end{bmatrix}}_{1 \times (d_v/2)} \tag{47}$$

to avoid notational clutter, we write θ^{t/d_v} as θ^t for short.

$$\text{freqs} = \underbrace{\begin{bmatrix} \frac{1}{\theta^0} & \frac{1}{\theta^2} & \frac{1}{\theta^4} & \cdots & \frac{1}{\theta^{d_v-2}} \end{bmatrix}}_{1 \times (d_v/2)} \tag{48}$$

this is in the decreasing frequency order.

```
        # Compute angles for all positions
        pos = torch.arange(x.shape[2])
        angles = torch.outer(pos, freqs)
```

since x shape is (batch, num_heads, T, head_dim), therefore x.shape[2] is the seq_len, i.e., T. Therefore, we have:

$$
\begin{aligned}
\text{angles} &= \begin{bmatrix}
0 \times \frac{1}{\theta^0} & 0 \times \frac{1}{\theta^2} & 0 \times \frac{1}{\theta^4} & \cdots & 0 \times \frac{1}{\theta^{d_v-2}} \\
1 \times \frac{1}{\theta^0} & 1 \times \frac{1}{\theta^2} & 1 \times \frac{1}{\theta^4} & \cdots & 1 \times \frac{1}{\theta^{d_v-2}} \\
\vdots & \vdots & \vdots & \ddots & \vdots \\
(T-1) \times \frac{1}{\theta^0} & (T-1) \times \frac{1}{\theta^2} & (T-1) \times \frac{1}{\theta^4} & \cdots & (T-1) \times \frac{1}{\theta^{d_v-2}}
\end{bmatrix} \\
&= \underbrace{\begin{bmatrix}
\frac{0}{\theta^0} & \frac{0}{\theta^2} & \frac{0}{\theta^4} & \cdots & \frac{0}{\theta^{d_v-2}} \\
\frac{1}{\theta^0} & \frac{1}{\theta^2} & \frac{1}{\theta^4} & \cdots & \frac{1}{\theta^{d_v-2}} \\
\vdots & \vdots & \vdots & \ddots & \vdots \\
\frac{T-1}{\theta^0} & \frac{T-1}{\theta^2} & \frac{T-1}{\theta^4} & \cdots & \frac{T-1}{\theta^{d_v-2}}
\end{bmatrix}}_{T \times (d_v/2)}
\end{aligned}
\tag{49}
$$

for each t^{th} token, there is a corresponding frequency for each feature, i.e., $\frac{t}{\theta^0}, \frac{t}{\theta^2}, \frac{t}{\theta^4}, \ldots, \frac{t}{\theta^{d_v-2}}$

```
        # Precompute sin/cos
        cos = torch.cos(angles)
        sin = torch.sin(angles)
        cos = cos.unsqueeze(0).unsqueeze(0)
        sin = sin.unsqueeze(0).unsqueeze(0)
```

$$\cos = \left[\left[\left[\begin{array}{ccccc} \cos(\frac{0}{\theta^0}) & \cos(\frac{0}{\theta^2}) & \cos(\frac{0}{\theta^4}) & \cdots & \cos(\frac{0}{\theta^{d_v-2}}) \\ \cos(\frac{1}{\theta^0}) & \cos(\frac{1}{\theta^2}) & \cos(\frac{1}{\theta^4}) & \cdots & \cos(\frac{1}{\theta^{d_v-2}}) \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \cos(\frac{T-1}{\theta^0}) & \cos(\frac{T-1}{\theta^2}) & \cos(\frac{T-1}{\theta^4}) & \cdots & \cos(\frac{T-1}{\theta^{d_v-2}}) \end{array}\right]\right]\right] \tag{50}$$

$$\sin = \left[\left[\left[\begin{array}{ccccc} \sin(\frac{0}{\theta^0}) & \sin(\frac{0}{\theta^2}) & \sin(\frac{0}{\theta^4}) & \cdots & \sin(\frac{0}{\theta^{d_v-2}}) \\ \sin(\frac{1}{\theta^0}) & \sin(\frac{1}{\theta^2}) & \sin(\frac{1}{\theta^4}) & \cdots & \sin(\frac{1}{\theta^{d_v-2}}) \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \sin(\frac{T-1}{\theta^0}) & \sin(\frac{T-1}{\theta^2}) & \sin(\frac{T-1}{\theta^4}) & \cdots & \sin(\frac{T-1}{\theta^{d_v-2}}) \end{array}\right]\right]\right] \tag{51}$$

```
# Split into pairs
x_even = x[..., 0::2]
x_odd = x[..., 1::2]
```

for example, if x is $\mathbf{Q}$, then we have:

$$\mathbf{Q}_{even} = \begin{bmatrix} q_{0,0} & q_{0,2} & q_{0,4} & \cdots & q_{0,d_v-2} \\ q_{1,0} & q_{1,2} & q_{1,4} & \cdots & q_{1,d_v-2} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ q_{T-1,0} & q_{T-1,2} & q_{T-1,4} & \cdots & q_{T-1,d_v-2} \end{bmatrix} \quad \mathbf{Q}_{odd} = \begin{bmatrix} q_{0,1} & q_{0,3} & q_{0,5} & \cdots & q_{0,d_v-1} \\ q_{1,1} & q_{1,3} & q_{1,5} & \cdots & q_{1,d_v-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ q_{T-1,1} & q_{T-1,3} & q_{T-1,5} & \cdots & q_{T-1,d_v-1} \end{bmatrix} \tag{52}$$

```
# Apply rotation (broadcasting cos/sin across batch and heads)
x_even_new = x_even * cos - x_odd * sin
x_odd_new = x_even * sin + x_odd * cos
```

$$\mathbf{Q}'_{even} =$$
$$\underbrace{\begin{bmatrix} q_{0,0}\cos(\frac{0}{\theta^0}) - q_{0,1}\sin(\frac{0}{\theta^0}) & q_{0,2}\cos(\frac{0}{\theta^2}) - q_{0,3}\sin(\frac{0}{\theta^2}) & \cdots & q_{0,d_v-2}\cos(\frac{0}{\theta^{d_v-2}}) - q_{0,d_v-1}\sin(\frac{0}{\theta^{d_v-2}}) \\ q_{1,0}\cos(\frac{1}{\theta^0}) - q_{1,1}\sin(\frac{1}{\theta^0}) & \textcolor{red}{q_{1,2}\cos(\frac{1}{\theta^2}) - q_{1,3}\sin(\frac{1}{\theta^2})} & \cdots & q_{1,d_v-2}\cos(\frac{1}{\theta^{d_v-2}}) - q_{1,d_v-1}\sin(\frac{1}{\theta^{d_v-2}}) \\ \vdots & \vdots & \vdots & \ddots \\ q_{T-1,0}\cos(\frac{T-1}{\theta^0}) - q_{T-1,1}\sin(\frac{T-1}{\theta^0}) & q_{T-1,2}\cos(\frac{T-1}{\theta^2}) - q_{T-1,3}\sin(\frac{T-1}{\theta^2}) & \cdots & q_{T-1,d_v-2}\cos(\frac{T-1}{\theta^{d_v-2}}) - q_{T-1,d_v-1}\sin(\frac{T-1}{\theta^{d_v-2}}) \\ q_{T,0}\cos(\frac{T}{\theta^0}) - q_{T,1}\sin(\frac{T}{\theta^0}) & q_{T,2}\cos(\frac{T}{\theta^2}) - q_{T,3}\sin(\frac{T}{\theta^2}) & \cdots & q_{T,d_v-2}\cos(\frac{T}{\theta^{d_v-2}}) - q_{T,d_v-1}\sin(\frac{T}{\theta^{d_v-2}}) \end{bmatrix}}_{T \times (d_v/2)}$$
$$\tag{53}$$

$$\mathbf{Q}'_{odd} =$$

$$\begin{bmatrix} q_{0,0}\sin(\frac{0}{\theta^0}) + q_{0,1}\cos(\frac{0}{\theta^0}) & q_{0,2}\sin(\frac{0}{\theta^2}) + q_{0,3}\cos(\frac{0}{\theta^2}) & \cdots & q_{0,d_v-2}\sin(\frac{0}{\theta^{d_v-2}}) + q_{0,d_v-1}\cos(\frac{0}{\theta^{d_v-2}}) \\[2mm] q_{1,0}\sin(\frac{1}{\theta^0}) + q_{1,1}\cos(\frac{1}{\theta^0}) & \color{red}{q_{1,2}\sin(\frac{1}{\theta^2}) + q_{1,3}\cos(\frac{1}{\theta^2})} & \cdots & q_{1,d_v-2}\sin(\frac{1}{\theta^{d_v-2}}) + q_{1,d_v-1}\cos(\frac{1}{\theta^{d_v-2}}) \\[2mm] \vdots & \vdots & \vdots & \ddots \\[2mm] q_{T-1,0}\sin(\frac{T-1}{\theta^0}) + q_{T-1,1}\cos(\frac{T-1}{\theta^0}) & q_{T-1,2}\sin(\frac{T-1}{\theta^2}) + q_{T-1,3}\cos(\frac{T-1}{\theta^2}) & \cdots & q_{T-1,d_v-2}\sin(\frac{T-1}{\theta^{d_v-2}}) + q_{T-1,d_v-1}\cos(\frac{T-1}{\theta^{d_v-2}}) \\[2mm] q_{T,0}\sin(\frac{T}{\theta^0}) + q_{T,1}\cos(\frac{T}{\theta^0}) & q_{T,2}\sin(\frac{T}{\theta^2}) + q_{T,3}\cos(\frac{T}{\theta^2}) & \cdots & q_{T,d_v-2}\sin(\frac{T}{\theta^{d_v-2}}) + q_{T,d_v-1}\cos(\frac{T}{\theta^{d_v-2}}) \end{bmatrix}$$

$$\underbrace{}_{T \times (d_v/2)}$$

$$(54)$$

One can see that the addition of the corresponding even and odd elements will be a $2-d$ rotation by each corresponding θ with different frequency. For example, the addtion of the two blue terms will be a $2-d$ rotation by θ^2 with frequency 1, i.e., by $\frac{1}{\theta^2}$ radians, and the addtion of the two red terms will be a $2-d$ rotation by $\frac{1}{\theta^2}$ radians.

$$\begin{bmatrix} q'_{1,2} \\ q'_{1,3} \end{bmatrix} = \begin{bmatrix} \cos(\frac{1}{\theta^2}) & -\sin(\frac{1}{\theta^2}) \\ \sin(\frac{1}{\theta^2}) & \cos(\frac{1}{\theta^2}) \end{bmatrix} \begin{bmatrix} q_{1,2} \\ q_{1,3} \end{bmatrix} = \begin{bmatrix} q_{1,2}\cos(\frac{1}{\theta^2}) - q_{1,3}\sin(\frac{1}{\theta^2}) \\ q_{1,2}\sin(\frac{1}{\theta^2}) + q_{1,3}\cos(\frac{1}{\theta^2}) \end{bmatrix} \tag{55}$$

or more generally, for each token t, and each feature block $(2i, 2i+1)$ (i.e., the i^{th} feature block), we have:

$$\begin{aligned} \begin{bmatrix} q'_{t,2i} & q'_{t,2i+1} \end{bmatrix} &= \begin{bmatrix} q_{t,2i}\cos(\frac{t}{\theta^2}) - q_{t,2i+1}\sin(\frac{t}{\theta^2}) & q_{t,2i}\sin(\frac{t}{\theta^2}) + q_{t,2i+1}\cos(\frac{t}{\theta^2}) \end{bmatrix} \\ &= \begin{bmatrix} q_{t,2i} & q_{t,2i+1} \end{bmatrix} R(\frac{t}{\theta^2}) \end{aligned} \tag{56}$$

```
# Interleave back
x_rotated = torch.stack([x_even_new, x_odd_new], dim=-1)
x_rotated = x_rotated.flatten(-2, -1)
return x_rotated
```

torch.stack([], dim=-1) will stack the $\mathbf{Q}'_{even}$, of size $(T, (d_v/2))$, and $\mathbf{Q}'_{odd}$, of size $(T, (d_v/2))$, along the last dimension (hence creating a new dimension of size 2) so that the result is a tensor of shape $(T, (d_v/2), 2)$, and then the flatten(-2, -1) will flatten the last two dimensions, so that the result is a tensor of shape (T, d_v), which is the same as the original $\mathbf{Q}$ tensor.

4.1 Decoupled RoPE

we have, given h heads:

$$\mathbf{Y}^{(i)} = \text{softmax}\left(\frac{1}{\sqrt{d_k}} \underbrace{\mathbf{X}}_{T \times d_{\text{model}}} \underbrace{\mathbf{W}_Q^{(i)}}_{d_{\text{model}} \times 128} \underbrace{\mathbf{W}_{\uparrow K}^{(i)\top}}_{128 \times 576} \underbrace{(\mathbf{X}\,\mathbf{W}_{\downarrow K,V})^\top}_{576 \times T} \right) \underbrace{\mathbf{X}\,\mathbf{W}_{\downarrow K,V}}_{T \times 576} \underbrace{\mathbf{W}_{\uparrow V}^{(i)}}_{576 \times d_v} \underbrace{\mathbf{W}_O^{(i-1) \times d_v + 1 : i \times d_v,\, 1 : d_{\text{model}}}}_{d_v \times d_{\text{model}}}$$

$$\mathbf{Y} = \sum_{i=1}^{h} \mathbf{Y}^{(i)}$$

(57)

when we have multiple heads, the output is the same size, except there are multiple of $\mathbf{Y}^{(i)}$ added together.

Here is something about RoPE, let's take the terms inside the softmax function, and remove the $\frac{1}{\sqrt{d_k}}$ term, we have:

$$\underbrace{\mathbf{X}}_{T \times d_{\text{model}}} \underbrace{\mathbf{W}_Q^{(i)}}_{d_{\text{model}} \times 128} \underbrace{\mathbf{W}_{\uparrow K}^{(i)\top}}_{128 \times 576} \underbrace{(\mathbf{X}\,\mathbf{W}_{\downarrow K,V})^\top}_{576 \times T}$$

(58)

Since RoPE is applied to the $\mathbf{Q}$ and $\mathbf{K}$ matrices, where there is not a single RoPE matrix for all the tokens, i.e., we do not have something like:

$$\underbrace{\mathbf{X}}_{T \times d_{\text{model}}} \underbrace{\mathbf{W}_Q^{(i)}}_{d_{\text{model}} \times 128} \mathbf{R}^{(i)} \mathbf{R}^{(i)T} \underbrace{\mathbf{W}_{\uparrow K}^{(i)\top}}_{128 \times 576} \underbrace{(\mathbf{X}\,\mathbf{W}_{\downarrow K,V})^\top}_{576 \times T}$$

(59)

For this reason, we may consider a decoupled RoPE mechanism, where we have two separate part of each row of $\mathbf{Q}$ and $\mathbf{K}$ matrices, i.e., we break up $\mathbf{q}_t$ and $\mathbf{k}_t$ into two parts, one is the part that does not involve RoPE (of a larger dimension), and the other is the part that involves RoPE (of a smaller dimension).

For each head i, and each token t, we have:

$$\boldsymbol{k}_t^C = \boldsymbol{W}_{\uparrow K}^{(i)} \, \boldsymbol{c}_t^{KV}, \quad \boldsymbol{v}_t^C = \boldsymbol{W}_{\uparrow V}^{(i)} \, \boldsymbol{c}_t^{KV}$$

(60)

for the t-th token, we have the following concatenation of outputs:

$$\boldsymbol{o}_t = \begin{bmatrix} \boldsymbol{o}_{t,1} & \boldsymbol{o}_{t,2} & \dots & \boldsymbol{o}_{t,h} \end{bmatrix}$$

(61)

which is the weighted sum of $\boldsymbol{v}_j^C$ for $j = 1, 2, \dots, t$, responding to a specific query $\boldsymbol{q}_t$, we now drop the head index i for simplicy:

$$\boldsymbol{o}_t = \sum_{j=1}^{t} \text{Softmax}\left(\frac{\boldsymbol{q}_t^\top \boldsymbol{k}_j}{\sqrt{d_h + d_h^R}} \right) \boldsymbol{v}_j^C$$

(62)

by letting: $\boldsymbol{c}_t^Q = \boldsymbol{W}_{\downarrow Q} \, \boldsymbol{h}_t, \quad \boldsymbol{c}_t^Q \in \mathbb{R}^{d_C}$, we have:

$$q_t = \left[(\underbrace{\boldsymbol{W}_{\uparrow Q}}_{d_h \times d_C} \ \underbrace{\boldsymbol{W}_{\downarrow Q}}_{d_C \times d_{\text{model}}} \ \boldsymbol{h}_t) \ \ \text{RoPE}(\underbrace{\boldsymbol{W}_Q^R}_{d_h^R \times d_C} \ \underbrace{\boldsymbol{W}_{\downarrow Q}}_{d_C \times d_{\text{model}}} \ \boldsymbol{h}_t) \right]$$

$$= \underbrace{\left[\boldsymbol{W}_{\uparrow Q} \ \boldsymbol{c}_t^Q \ \ \text{RoPE}(\boldsymbol{W}_Q^R \ \boldsymbol{c}_t^Q) \right]}_{(d_h + d_h^R) \times 1} \tag{63}$$

you see that $\mathbf{q}_t$ has two parts, each projecting the hidden state $\mathbf{h}_t$ into a different space by matrices $\boldsymbol{W}_{\uparrow Q}$ and $\boldsymbol{W}_Q^R$. Similiarly, by letting: $\boldsymbol{c}_t^{KV} = \boldsymbol{W}_{\downarrow K} \ \boldsymbol{h}_t, \quad \boldsymbol{c}_t^{KV} \in \mathbb{R}^{d_C}$, we have:

$$k_j = \left[(\underbrace{\boldsymbol{W}_{\uparrow K}}_{d_h \times d_C} \ \underbrace{\boldsymbol{W}_{\downarrow K}}_{d_C \times d_{\text{model}}} \ \boldsymbol{h}_j) \ \ \text{RoPE}(\underbrace{\boldsymbol{W}_K^R}_{d_h^R \times d_C} \ \underbrace{\boldsymbol{W}_{\downarrow K}}_{d_C \times d_{\text{model}}} \ \boldsymbol{h}_j) \right]$$

$$= \underbrace{\left[\boldsymbol{W}_{\uparrow K} \ \boldsymbol{c}_j^{KV} \ \ \text{RoPE}(\boldsymbol{W}_K^R \ \boldsymbol{c}_j^{KV}) \right]}_{(d_h + d_h^R) \times 1} \tag{64}$$

$$k_t = \left[\boldsymbol{W}_{\uparrow K} \ \boldsymbol{c}_t^{KV} \ \ \text{RoPE}(\boldsymbol{W}_K^R \ \boldsymbol{c}_t^{KV}) \right] \in \mathbb{R}^{d_h + d_h^R}$$

$$\text{where} \quad \boldsymbol{W}_{\uparrow K} \in \mathbb{R}^{d_h \times d_C}, \quad \boldsymbol{W}_K^R \in \mathbb{R}^{d_h^R \times d_C} \tag{65}$$

now we look at the value $\boldsymbol{v}_t^C$, we have:

$$v_t^C = \underbrace{\boldsymbol{W}_{\uparrow V}}_{d_h \times d_C} \ \underbrace{\boldsymbol{W}_{\downarrow KV}}_{d_C \times d_{\text{model}}} \ \underbrace{\boldsymbol{h}_t}_{d_{\text{model}} \times 1}$$

$$= \underbrace{\boldsymbol{W}_{\uparrow V}}_{d_h \times d_C} \ \underbrace{\boldsymbol{c}_t^{KV}}_{d_C \times 1} \tag{66}$$

during inference, we have the following steps:

- for the t-th token, we compute:

$$q_t = \left[\boldsymbol{W}_{\uparrow Q} \ \boldsymbol{c}_t^Q \ \ \text{RoPE}(\boldsymbol{W}_Q^R \ \boldsymbol{c}_t^Q) \right] \tag{67}$$

here, nothing needs to be cached.

- and for each of the j-th token, $\forall j \in \{1, 2, \ldots, t-1\}$, we have:

$$k_j = \left[\boldsymbol{W}_{\uparrow K} \ \boldsymbol{c}_j^{KV} \ \ \text{RoPE}(\boldsymbol{W}_K^R \ \boldsymbol{c}_j^{KV}) \right] \tag{68}$$

the things in the cache are:

  - $\boldsymbol{c}_t^{KV} = \boldsymbol{W}_{\downarrow KV} \ \boldsymbol{h}_t$
  - $\boldsymbol{k}_t^R = \text{RoPE}(\boldsymbol{W}_K^R \ \boldsymbol{c}_t^{KV})$

5 Flash Attention

5.1 traditional softmax algorithm

The algorithm for computing softmax is as follows, given there is a logit set $\{x_i\}_{i=1}^N$, where each $x_i \in \mathbb{R}$. We need three separate loops to compute the softmax output:

1. loop 1: compute the maximum value in the logit set:

$$
\begin{aligned}
&m_0 = -\infty \\
&\text{for } 1 \leq i \leq N \text{ do} \\
&\quad m_i = \max(m_{i-1}, x_i) \\
&\text{return } m_N
\end{aligned}
\tag{69}
$$

2. loop 2: compute the sum of the exponential of the logits after we obtained m_N which is the maximum value in the logit set:

$$
\begin{aligned}
&d_0 = 0 \\
&\text{for } 1 \leq i \leq N \text{ do} \\
&\quad d_i = d_{i-1} + \exp^{x_i - m_N} \\
&\text{return } d_N
\end{aligned}
\tag{70}
$$

3. loop 3: normalize the logits:

$$
\begin{aligned}
&\text{for } 1 \leq i \leq N \text{ do} \\
&\quad a_i = \frac{\exp^{x_i - m_N}}{d_N} \\
&\text{return } \{a_i\}_{i=1}^N
\end{aligned}
\tag{71}
$$

5.2 online softmax algorithm

now looking at the second loop, in essence, at the i-th iteration, it is computing the sum of the exponential of the logits from 1 to i:

$$
\begin{aligned}
d_i &= \sum_{j=1}^{i} \exp^{x_j - m_N} \\
&= \underbrace{\sum_{j=1}^{i-1} \exp^{x_j - m_N}}_{d_{i-1}} + \exp^{x_i - m_N} \\
&= d_{i-1} + \exp^{x_i - m_N}
\end{aligned}
\tag{72}
$$

with the initial condition $d_0 = 0$, the above can also be expressed as:

22

$$d_0 = 0 \quad \text{and} \quad d_i = d_{i-1} + \exp^{x_i - m_N} \tag{73}$$

However, the above cannot be combined with the first loop, because of the m_N term. Therefore, instead of subtracting a "global" maximum value m_N, we subtract a "intermediate" maximum value m_i up to the i-th iteration:

$$d'_i = \sum_{j=1}^{i} \exp^{x_j - m_i} \tag{74}$$

at the completion of the entire loop, we have $d'_N = d_N$, even though $d_i \neq d'_i \quad \forall i \neq N$ in general. However, if we do the same trick as Eq.(74), however, the first term is not equal to $d'_{i-1} = \sum_{j=1}^{i-1} \exp^{x_j - m_{i-1}}$, i.e., you cannot write the sum as, (a term only contain sum up to $i-1$ terms d_{i-1}) + (a term only contain x_i term).

$$d'_i = \underbrace{\sum_{j=1}^{i-1} \exp^{x_j - m_i}}_{\neq d'_{i-1}} + \exp^{x_i - m_i} \tag{75}$$

but we can make it so, by adding and subtracting a dummy variable:

$$
\begin{aligned}
d'_i &= \sum_{j=1}^{i-1} \exp^{x_j - m_i} + \exp^{x_i - m_i} \\
&= \Big(\sum_{j=1}^{i-1} \exp^{x_j - m_i} \Big) \underbrace{\exp^{m_{i-1}} \exp^{-m_{i-1}}}_{\text{dummy variable}} + \exp^{x_i - m_i} \quad \text{perform} \times \exp^{m_{i-1}} \exp^{-m_{i-1}} \\
&= \Big(\sum_{j=1}^{i-1} \exp^{x_j - m_{i-1}} \Big) \exp^{m_{i-1}} \exp^{-m_i} + \exp^{x_i - m_i} \quad \text{sawp the red and blue terms} \\
&= \underbrace{\Big(\sum_{j=1}^{i-1} \exp^{x_j - m_{i-1}} \Big)}_{d'_{i-1}} \exp^{m_{i-1} - m_i} + \exp^{x_i - m_i} \\
&= d'_{i-1} \exp^{m_{i-1} - m_i} + \exp^{x_i - m_i}
\end{aligned}
\tag{76}
$$

therefore, we can have the algorithm to be:

1. loop 1: compute the maximum value in the logit set:

$$\begin{aligned}
&m_0 = -\infty \\
&d_0 = 0 \\
&\text{for } 1 \leq i \leq N \text{ do} \\
&\quad m_i = \max(m_{i-1}, x_i) \\
&\quad d_i = d_{i-1}\exp^{m_{i-1}-m_i} + \exp^{x_i - m_i} \\
&\text{return } m_N, d_N
\end{aligned} \tag{77}$$

we also changed $d_i' \to d_i$ to simplify the notation.

2. loop 2: same as loop 3 in the traditional softmax algorithm:

$$\begin{aligned}
&\text{for } 1 \leq i \leq N \text{ do} \\
&\quad a_i = \frac{\exp^{x_i - m_N}}{d_N} \\
&\text{return } \{a_i\}_{i=1}^N
\end{aligned} \tag{78}$$

since we are not perform "reduce" operation, such as sum, max, etc., as in here, we compute each a_i individually, we cannot combine it with the first loop.

5.3   online attention

looking at the single $\mathbf{q}$ row matrix form, where we have:

$$\begin{aligned}
\mathbf{q}\mathbf{K}^\top &= \begin{bmatrix} \mathbf{q}\mathbf{k}_1^\top & \dots & \mathbf{q}\mathbf{k}_m^\top \end{bmatrix} \\
\implies \mathbf{o} &= \operatorname{softmax}(\mathbf{q}\mathbf{K}^\top)\mathbf{V} \\
&= \begin{bmatrix} \operatorname{softmax}\left(\begin{bmatrix} \mathbf{q}\mathbf{k}_1^\top & \dots & \mathbf{q}\mathbf{k}_m^\top \end{bmatrix} \right) \end{bmatrix} \begin{bmatrix} - & \mathbf{v}_1 & - \\ - & \dots & - \\ - & \mathbf{v}_m & - \end{bmatrix} \\
&= \sum_{i=1}^N \frac{\exp[\mathbf{q}\mathbf{k}_i^\top]}{\sum_j \exp[\mathbf{q}\mathbf{k}_j^\top]}\mathbf{v}_i \\
&= \sum_{i=1}^N a_i \mathbf{v}_i
\end{aligned} \tag{79}$$

taking a particular row of the above, $\mathbf{x} = \mathbf{q}\mathbf{K}^\top$, where $x_i = \mathbf{q}\mathbf{k}_i^\top$ is an element of $\mathbf{x}$:

1. loop 1: compute the maximum value in the logit set:

$$m_0 = -\infty$$
$$d_0 = 0$$
$$\text{for } 1 \leq i \leq N \text{ do}$$
$$\quad \color{red}{x_i = q_i \mathbf{k}_i^\top}$$
$$\quad m_i = \max(m_{i-1}, x_i)$$
$$\quad d_i = d_{i-1} \exp^{m_{i-1}-m_i} + \exp^{x_i-m_i}$$
$$\text{return } m_N, d_N$$

(80)

2. loop 2: adding $\mathbf{o} = \sum_{i=1}^{N} a_i \mathbf{v}_i$ to the second loop:

$$\mathbf{o}_0 = \mathbf{0}$$
$$\text{for } 1 \leq i \leq N \text{ do}$$
$$\quad a_i = \frac{\exp^{x_i-m_N}}{d_N}$$
$$\quad \mathbf{o}_i = \mathbf{o}_{i-1} + a_i \mathbf{v}_i$$

$$\mathbf{o}_0 = \mathbf{0}$$
$$\text{for } 1 \leq i \leq N \text{ do}$$
$$\quad \mathbf{o}_i = \mathbf{o}_{i-1} + \frac{\exp^{x_i-m_N}}{d_N} \mathbf{v}_i$$

(81)

note that we can apply the same trick as Eq.(74) to the second loop:

$$\mathbf{o}_i = \sum_{i=1}^{N} \frac{\exp^{x_i-m_N}}{d_N'} \mathbf{v}_i$$

$$\mathbf{o}_i' = \sum_{j=1}^{i} \frac{\exp^{x_j-m_i}}{d_i'} \mathbf{v}_j$$

(82)

where we have $\mathbf{o}_N = \mathbf{o}_N'$. Apply the same track as Eq.(74), we have:

$$
\begin{aligned}
\mathbf{o}_i' &= \sum_{j=1}^{i} \frac{\exp^{x_j-m_i}}{d_i'} \mathbf{v}_j \\
&= \Big(\sum_{j=1}^{i-1} \frac{\exp^{x_j-m_i}}{d_i'} \mathbf{v}_j \Big) + \frac{\exp^{x_i-m_i}}{d_i'} \mathbf{v}_i \\
&= \underbrace{\Big(\sum_{j=1}^{i-1} \frac{\exp^{x_j-m_i}}{d_i'} \mathbf{v}_j \Big) \frac{d_{i-1}'}{d_{i-1}'} \frac{\exp^{m_{i-1}}}{\exp^{m_{i-1}}}}_{\neq \mathbf{o}_{i-1}} + \frac{\exp^{x_i-m_i}}{d_i'} \mathbf{v}_i \\
&= \underbrace{\Big(\sum_{j=1}^{i-1} \frac{\exp^{x_j-m_{i-1}}}{d_{i-1}'} \mathbf{v}_j \Big)}_{= \mathbf{o}_{i-1}} \frac{d_{i-1}'}{d_i'} \frac{\exp^{m_i}}{\exp^{m_{i-1}}} + \frac{\exp^{x_i-m_i}}{d_i'} \mathbf{v}_i \\
&= \mathbf{o}_{i-1} \frac{d_{i-1}'}{d_i'} \frac{\exp^{m_i}}{\exp^{m_{i-1}}} + \frac{\exp^{x_i-m_i}}{d_i'} \mathbf{v}_i
\end{aligned}
$$

(83)

5.3.1 final algorithm

everything comes down to just a single loop:

$$
\begin{aligned}
&m_0 = -\infty \\
&d_0 = 0 \\
&\mathbf{o}_0 = \mathbf{0} \\
&\text{for } 1 \leq i \leq N \text{ do} \\
&\qquad x_i = q_i \mathbf{k}_i^\top \\
&\qquad m_i = \max(m_{i-1}, x_i) \\
&\qquad d_i = d_{i-1} \exp^{m_{i-1} - m_i} + \exp^{x_i - m_i} \\
&\qquad \mathbf{o}_i = \mathbf{o}_{i-1} \frac{d_{i-1}}{d_i} \frac{\exp^{m_i}}{\exp^{m_{i-1}}} + \frac{\exp^{x_i - m_i}}{d_i} \mathbf{v}_i \\
&\text{return } \mathbf{o}_N
\end{aligned}
\tag{84}
$$

again, we changed $d_i' \to d_i$ and $\mathbf{o}_i' \to \mathbf{o}_i$ to simplify the notation.

References

[1] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Lukasz Kaiser, and Illia Polosukhin, "Attention is all you need," Advances in neural information processing systems, vol. 30, 2017.